

# modeling-for-business-forecasting

February 7, 2024

## Data Collection and Preprocessing:

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import PolynomialFeatures
from sklearn.impute import SimpleImputer
```

```
[ ]: data = pd.read_csv("/content/Financial Distress.csv")
data.head()
```

```
[ ]: Company Time Financial Distress x1 x2 x3 x4 \
0 1 1 0.010636 1.2810 0.022934 0.87454 1.21640
1 1 2 -0.455970 1.2700 0.006454 0.82067 1.00490
2 1 3 -0.325390 1.0529 -0.059379 0.92242 0.72926
3 1 4 -0.566570 1.1131 -0.015229 0.85888 0.80974
4 2 1 1.357300 1.0623 0.107020 0.81460 0.83593
```

```
 x5 x6 x7 ... x74 x75 x76 x77 x78 \
0 0.060940 0.188270 0.52510 ... 85.437 27.07 26.102 16.000 16.0
1 -0.014080 0.181040 0.62288 ... 107.090 31.31 30.194 17.000 16.0
2 0.020476 0.044865 0.43292 ... 120.870 36.07 35.273 17.000 15.0
3 0.076037 0.091033 0.67546 ... 54.806 39.80 38.377 17.167 16.0
4 0.199960 0.047800 0.74200 ... 85.437 27.07 26.102 16.000 16.0
```

```
 x79 x80 x81 x82 x83
0 0.2 22 0.060390 30 49
1 0.4 22 0.010636 31 50
2 -0.2 22 -0.455970 32 51
3 5.6 22 -0.325390 33 52
4 0.2 29 1.251000 7 27
```

[5 rows x 86 columns]

```
[ ]: print(data.isnull().sum())
```

```
Company      0
Time         0
```

```
Financial Distress    0
x1                    0
x2                    0
..
x79                   0
x80                   0
x81                   0
x82                   0
x83                   0
Length: 86, dtype: int64
```

### Feature Selection and Engineering:

```
[ ]: #correlation_matrix
correlation_matrix = data.corr().abs()
upper_tri = correlation_matrix.where(np.triu(np.ones(correlation_matrix.shape),
↪k=1).astype(np.bool))
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.
↪95)]
data.drop(to_drop, axis=1, inplace=True)
```

<ipython-input-19-891450e28ae9>:3: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool\_` here.  
 Deprecated in NumPy 1.20; for more details and guidance:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
upper_tri =
correlation_matrix.where(np.triu(np.ones(correlation_matrix.shape),
k=1).astype(np.bool))
```

```
[ ]: # Define the features (X) and target variable (y)
X = data.drop(columns=['Financial Distress'])
y = data['Financial Distress']
```

```
[ ]: # Handle missing values by imputing them with the mean of each column
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)

# Feature engineering: Creating new features that might improve the model's
↪predictive power
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly.fit_transform(X_imputed)
```

```
[ ]: data['month'] = pd.to_datetime(data['Time']).dt.month
data['day'] = pd.to_datetime(data['Time']).dt.day
data['hour'] = pd.to_datetime(data['Time']).dt.hour
```

```
[ ]: data['x1_lag1'] = data['x1'].shift(1)
```

### Model Selection:

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.impute import SimpleImputer
```

```
[ ]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[ ]: # Handle missing values in the features
imputer = SimpleImputer(strategy='mean')
X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

```
[ ]: # Model selection: Choose appropriate regression models
models = {
    "Linear Regression": LinearRegression(),
    "Random Forest Regressor": RandomForestRegressor(random_state=42)
}
```

```
[ ]: # Train and evaluate each model
for name, model in models.items():
    model.fit(X_train_imputed, y_train)
    y_pred = model.predict(X_test_imputed)
    mse = mean_squared_error(y_test, y_pred)
    print(f"{name}: Mean Squared Error = {mse:.2f}")
```

Linear Regression: Mean Squared Error = 1003.63

Random Forest Regressor: Mean Squared Error = 1.14

### Model Development and Evaluation:

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error
```

```
[ ]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[ ]: # Handle missing values in the features
imputer = SimpleImputer(strategy='mean')
```

```
X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

```
[ ]: # Define the model
model = RandomForestRegressor(random_state=42)
```

```
[ ]: # Train the model
model.fit(X_train_imputed, y_train)
```

```
[ ]: RandomForestRegressor(random_state=42)
```

```
[ ]: # Make predictions on the testing data
y_pred = model.predict(X_test_imputed)
```

```
[ ]: # Evaluate the model using Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse:.2f}")
```

Mean Squared Error: 1.14

### Predictive Insights:

```
[ ]: # Get feature importances
feature_importance = model.feature_importances_

# Sort feature importances in descending order
sorted_idx = np.argsort(feature_importance)[::-1]

# Print feature names and their importance scores
for i in sorted_idx:
    print(f"{X.columns[i]}: {feature_importance[i]}")
```

```
x48: 0.3090087194062596
x41: 0.163136481522076
x25: 0.14561107025407827
x10: 0.04382329461971983
x73: 0.017246554441904243
x19: 0.015594210974306778
x9: 0.015144994870977339
x83: 0.010487358954927717
x3: 0.00891042212655669
x28: 0.008371355916789953
x14: 0.007770702679828272
x47: 0.007670687128422793
x30: 0.007654078740588396
x58: 0.007427739314997659
x39: 0.007114408382947107
x59: 0.007058183524293813
```

x22: 0.006987261700509246  
x21: 0.0068025684156600204  
x32: 0.006608977269191928  
x68: 0.006439945229182956  
x20: 0.0062812419490220605  
x24: 0.0061139003514323  
x54: 0.006013698947933399  
x36: 0.005988819716260606  
x46: 0.005668338677229048  
x66: 0.005658450546249864  
x16: 0.005499047953228381  
x40: 0.00527019042111605  
x12: 0.005175151527174476  
x45: 0.005141277201464576  
x18: 0.005139162472876233  
x71: 0.004953625676704989  
x61: 0.004786823114997776  
x55: 0.004723718066318178  
x37: 0.004660609151906636  
x57: 0.004633516060501674  
x79: 0.004587664183173071  
x63: 0.004247660829974798  
x5: 0.003946581653934656  
x53: 0.003924205272744166  
x17: 0.003787997113625004  
x31: 0.0037141194773741533  
x74: 0.0036646564860900693  
x65: 0.0036126862164825076  
x27: 0.0033884080396323845  
x60: 0.0032770944007498855  
x56: 0.0031738224919303764  
x38: 0.003155396543887233  
Company: 0.0029481489006722684  
x33: 0.0028328049504060334  
x51: 0.002802768096463701  
x23: 0.0026232040416876885  
x26: 0.0025735433340151738  
x62: 0.002549688068575275  
x11: 0.0025192702926476822  
x43: 0.0025185497287355464  
x64: 0.0025171437207263335  
x80: 0.002496805505393409  
x44: 0.0024832000490814256  
x4: 0.0024805509735338306  
x82: 0.0024366758203886777  
x2: 0.0024125354015970017  
x69: 0.00235748688058591  
x8: 0.0022836939218970984

```

x70: 0.002037548427627404
x15: 0.0020169254097200874
x67: 0.0020091791483752737
Time: 0.0019642868670897273
x42: 0.0019554510512905495
x7: 0.001909992801898369
x35: 0.0016277879489805971
x6: 0.0014914246773651393
x29: 0.0014912068788452888
x1: 0.0014718832262057494
x72: 0.000851730491901645
x78: 0.0007274716320505119
x77: 0.0005521617350094228

```

```

[ ]: # Calculate residuals
residuals = y_test - y_pred

```

```

[ ]: # Find instances with high prediction errors
high_error_indices = np.argsort(np.abs(residuals))[:, :-1][:10] # Adjust the
↳ number of instances as needed

# Print feature values and their corresponding actual and predicted values for
↳ high-error instances
for idx in high_error_indices:
    print(f"Instance {idx}:")
    print(f"Actual Value: {y_test.iloc[idx]}")
    print(f"Predicted Value: {y_pred[idx]}")
    print(f"Features: {X_test.iloc[idx]}")

```

```

Instance 347:
Actual Value: 9.3749
Predicted Value: 1.8606840999999992
Features: Company    96.00000
Time              4.00000
x1                0.70861
x2                0.58095
x3                0.84753
...
x78               16.00000
x79               5.60000
x80               18.00000
x82               4.00000
x83               15.00000
Name: 1073, Length: 77, dtype: float64
Instance 665:
Actual Value: -5.6838
Predicted Value: 0.6850457449999996
Features: Company    89.000000

```

Time 1.000000  
 x1 1.125600  
 x2 0.028183  
 x3 0.860420  
 ...  
 x78 16.000000  
 x79 0.200000  
 x80 23.000000  
 x82 3.000000  
 x83 20.000000  
 Name: 1006, Length: 77, dtype: float64  
 Instance 154:  
 Actual Value: 4.7789  
 Predicted Value: 11.0579138  
 Features: Company 284.00000  
 Time 11.00000  
 x1 6.84810  
 x2 0.74941  
 x3 0.20403  
 ...  
 x78 14.50000  
 x79 -13.75000  
 x80 25.00000  
 x82 2.00000  
 x83 9.00000  
 Name: 2893, Length: 77, dtype: float64  
 Instance 128:  
 Actual Value: 0.96119  
 Predicted Value: 6.404533899999997  
 Features: Company 360.000000  
 Time 9.000000  
 x1 0.726020  
 x2 0.068228  
 x3 0.787430  
 ...  
 x78 12.750000  
 x79 0.750000  
 x80 29.000000  
 x82 16.000000  
 x83 50.000000  
 Name: 3261, Length: 77, dtype: float64  
 Instance 240:  
 Actual Value: -2.7031  
 Predicted Value: 2.4237075520000007  
 Features: Company 15.000000  
 Time 6.000000  
 x1 0.989490  
 x2 0.080388

```

x3          0.762280
...
x78         12.000000
x79         -6.400000
x80         18.000000
x82         12.000000
x83         41.000000
Name: 139, Length: 77, dtype: float64
Instance 676:
Actual Value: 0.71129
Predicted Value: 5.7441793999999999
Features: Company    249.00000
Time              11.00000
x1                1.22220
x2                0.18870
x3                0.51988
...
x78              14.50000
x79             -16.00000
x80              19.00000
x82              16.00000
x83              55.00000
Name: 2599, Length: 77, dtype: float64
Instance 492:
Actual Value: 12.772
Predicted Value: 8.1241060999999995
Features: Company    284.00000
Time              10.00000
x1                1.28090
x2                0.52497
x3                0.38903
...
x78              14.12500
x79              -5.10000
x80              25.00000
x82              1.00000
x83              8.00000
Name: 2892, Length: 77, dtype: float64
Instance 545:
Actual Value: -0.72071
Predicted Value: 3.7704444399999995
Features: Company    263.00000
Time              7.00000
x1                1.11090
x2                0.11228
x3                0.83831
...
x78              12.00000

```



```

x79      -13.40000
x80       22.00000
x82        6.00000
x83       18.00000
Name: 2742, Length: 77, dtype: float64
Instance 46:
Actual Value: 1.5174
Predicted Value: 5.888393099999999
Features: Company    294.00000
Time        14.00000
x1          0.41990
x2          0.13382
x3          0.68937
...
x78       20.50000
x79        8.60000
x80       16.00000
x82        1.00000
x83        8.00000
Name: 2917, Length: 77, dtype: float64
Instance 399:
Actual Value: 1.7354
Predicted Value: 6.063268899999998
Features: Company    186.00000
Time        11.00000
x1          2.07540
x2          0.51572
x3          0.24060
...
x78       14.50000
x79      -16.00000
x80        4.00000
x82        1.00000
x83       17.00000
Name: 1992, Length: 77, dtype: float64

```

```

[ ]: # Plot residuals
plt.scatter(y_pred, residuals)
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residual Analysis")
plt.show()

```

